



Eidgenössische Technische Hochschule Zürich
Swiss Federal Institute of Technology Zurich

Lecture with Computer Exercises: Modelling and Simulating Social Systems with MATLAB

Project Report

Insert Title Here
...

Name 1 & Name 2

Zurich
May 2008

Agreement for free-download

We hereby agree to make our source code for this project freely available for download from the web pages of the SOMS chair. Furthermore, we assure that all source code is written by ourselves and is not violating any copyright restrictions.

Name 1

Name 2

Contents

1	Abstract	4
2	Individual contributions	4
3	Introduction and Motivations	4
3.1	What is Self-Organized Criticality?	4
3.2	What is Self-Organized Criticality	4
3.3	Flicker Noise	5
3.4	Forest-Fire Model	5
3.5	Power Laws	5
3.6	Key Parameters in SOC	6
3.7	Motivation	6
4	Description of the Model	7
4.1	Cellular automata	7
5	Implementation	7
5.1	Grid Based Updating	8
5.2	Random Based Updating	9
5.2.1	keeping track of the clusters	10
5.3	The <i>neighbor</i> function	10
5.4	Saving the data	11
6	Simulation Results and Discussion	11
6.1	Tree population	11
6.2	Cluster Radius	12
6.3	Cluster Size	13
7	Open Questions	14
8	Summary and Outlook	14

1 Abstract

2 Individual contributions

3 Introduction and Motivations

3.1 What is Self-Organized Criticality?

Self-Organized Criticality is the Concept that the dynamics of a physical system converge to a critical point independant of the bestowed boundary conditions. This critical point commonly is not the equilibrium point in a traditional, physical sense. The concept is best understood by example, so we shall bring up a very educative one introduced first by [2]). Imagine a sand pile in 2 dimensions. One can discretisze this pile into cells, which are denoted with the indices n . Each of those cells has an amount of sand grains on it, which is directly corresponding to the height z_n . Now we state the simple rule, that, if the difference in height between two neighboring cells becomes larger than a predefined limit, the higher cell will give grains to the lower cell (avalanche). Now every timestep, we drop a grain on cell one. Independant of the initial conditions, eventually, a straight slope will evolve. The critical point is reached if every difference between two neighboring cells is exactly the limit. What happens if we drop the next grain? Then the difference $z_1 - z_2 > p$ and one grain will drop on z_2 . Subsequently, the difference $z_2 - z_3$ will be too big and the process repeats until it hits the last cell. We therefore call a state critical if a yet so small disturbance can propagate throughout the whole system. A further condition for SOC is the presence of the critical point, in the example, that would be the slope that should form in every experiment independant of the starting conditions. It is important to note here, that this critical state is *not* the equilibrium state of the system, since that would be a flat surface.

3.2 What is Self-Organized Criticality

Self-organized criticality (SOC) is a concept applied to spacially extended dynamic physical systems. It is often quite difficult to describe the temporal and spatial evolution of such a system and a lot of research happens in that field. Usually, one uses a mean-field approach to such a problem, where the individual coupled degrees of freedom are replaced by a field. This (external) field then acts on the DOF. SOC is a different approach to this problem. It is the concept, that a certain class of dissipative, coupled dynamic systems converge to a critical point independantly of their underlying boundary and initial conditions. What are the properties of such a critical point? First of all, it is stable with respect to small disturbances. If this was

not true, the system would not evolve to such a point. Note that this does not imply that such a point *exists*. If a certain point is unstable, the system commonly will not evolve to this point. The critical point, however, generally is not the physical equilibrium point of the system.

3.3 Flicker Noise

Flicker noise is something often mentioned when talking about dissipative coupled dynamical systems. Flicker Noise (or $1/f$ -Noise) is the description for an often observed phenomenon of such systems; The observation, that the power spectrum $S(f)$ scales with $1/f$ or more generally, with $f^{-\beta}$. This implies that $S(f) \cdot f \propto 1$

3.4 Forest-Fire Model

A Forest Fire Model is basically a cellular automata. Each cell has three possible states:

1. Empty Site
2. Alive Tree
3. Burning Tree

There are four rules defined for each cell, which are executed simultaneously:

- A burning tree turns into an empty site
- A tree starts to burn if at least one of its neighbors are burning
- A tree will grow on an empty site randomly with probability p
- A tree will burn randomly with probability f

It is important to note that the Forest Fire Model does not only apply to forest fires as its name would suggest. A whole class of problems follows the same basic rules and can be treated accordingly. An example would be the spreading of a disease, where one can directly transform the above rules and states.

3.5 Power Laws

Power laws are functions of the form $P(s) \propto s^{-\tau}$. These functions are found in many data sets, such as the size distribution of cities in a certain area. Take any country, one will find one very large metropolis, a few large cities, many medium

sized towns and a huge lot of small towns. Power laws imply that the frequency of a general data point is inversely proportional to its magnitude.

In context to the FFM, it is proposed as a way of quantifying the frequency of Fires with respect to their cluster size.

Power laws have also prominently been found in data sets of earthquakes, where they apply almost perfectly, meaning that small earthquakes happen very often in respect to the frequency of big ones.

3.6 Key Parameters in SOC

To find out whether a basic FFM exhibits SOC behavior, we need to define a set of parameters which serves as indicators for the desired analysis.

- The time needed to burn down a forest cluster: This is heavily dependant on the form of the implementation. If we choose to implement it in a way like [reference needed], we will have an instantaneous fire which essentially takes no time to burn and just resets all the connected trees. If however, we choose the form of a visible fire, meaning that the ratio $f/p \ll 1$ is not 0 in limit and that it takes the fire one timestep to advance a grid cell, then the time needed to burn a forest cluster is an important variable we need to extract from the simulation.
- Number distribution of the size of clusters: Here we first need to define what a cluster is. A cluster is a set of neighboring cells all obtaining the same state. When we talk about forest clusters, we of course mean connected trees. In the implementation where the fire spreads infinitely fast, the forest cluster containing the ignitor cell is the same as the burnt area. The size distribution therefore is a very good indicator for the fire size distribution.
- The mean number of forest clusters in a unit volume $n(s)$, where s is the number of trees in a cluster.
- Number of fires per unit time step: This is a tuning parameter of the model.
- Correlation length.

We have tried to implement the model in a way that the measurement of these desired variables is possible with the least needed effort.

3.7 Motivation

Why is it important to study these effects? As mentioned before, forest fire models apply to a wide range of problems and are therefore helpful in predicting the behaviour of such systems. Self-Organized criticality is the concept that describes the

dynamics of such models and the two fields are therefore closely connected. One has closely studied these effects to make better predictions about forest fires. The main problem was, that in most cases, even relatively small fires were extinguished by the authorities. Since that led to an overcritical state, the probability for a huge, devastating fire rose quickly. This has happened in the past and had a serious impact on the ecosystem of those areas. Since the problem of self-organized criticality has been understood, one has stopped to extinguish every little fire, therefore avoiding to reach an overcritical state at which the disturbance (a small fire) can propagate throughout the whole system (large-scale fire). The same ideas can be applied to a variety of problems such as disease spreading or urban planning.

4 Description of the Model

4.1 Cellular automata

A cellular automata is a very powerful way of simulating problems which are defined by a set of rules. It is usually simulated on a grid, but not necessarily restricted to those geometrical constraints. The main characteristics are:

- Every grid point has a state
- Grid points change their state depending on the neighbor states according to a set of rules
- Random actions may be introduced

In the example of the FFM, every grid point has three states (empty, alive and burning) and four rules, the most obvious of which is change state to burning if you are alive and your neighbor is burning.

Cellular Automata can exhibit very complicated behavior when fed with very simple rules, which is the main reason why they are so powerful. There are similarities to agent-based models and networks, but it is not to be mistaken for one of these.

The second approach to grid-based updating is with a second grid that stores temporary information. [Insert Description here]

5 Implementation

As discussed before, there are two main implementation methods for the model: One way is to update the whole grid every timestep, whereas the second method picks a cell randomly at every timestep and only updates the chosen one. Both versions store the Grid in a Matrix which is initialized as follows:

```
Grid=zeros(Grid_Size);
```

or if the grid should be filled by a degree of k :

```
Grid=floor(rand(Grid_Size)+k);
```

5.1 Grid Based Updating

The basic construct of the grid based updating implementation looks like this:

```
for i=1:t % Loop over all timesteps
    for j=1:size(Grid,1) % Loop over all y-coordinates
        for k=1:size(Grid,2) % Loop over all x-coordinates
            if Grid(j,k)==0 %If the grid point is empty
                //Grid(j,k)=1 with some probability p
            end
        end
    end
    for j=1:size(Grid,1) % Loop over all y-coordinates
        for k=1:size(Grid,2) % Loop over all x-coordinates
            //if Neighbor(j,k)==3 % If grid neighbor is burning
            Grid(j,k)=2 % Set current grid point on fire
        end
    end
    for j=1:size(Grid,1) % Loop over all y-coordinates
        for k=1:size(Grid,2) % Loop over all x-coordinates
            if Grid(j,k)==3 % If Grid point is burning
                Grid(j,k)=0; % Turn it into an empty site
            end
            if Grid(j,k)==2 % If grid point is ignited
                Grid(j,k)=3; % Turn it into an empty site.
            end
        end
    end
end
end
```

Several things are to note here:

- Commands with `//` in front are not implemented exactly like that. Check the appendix for the complete source code.

- We need several spatial loops in order to successfully suppress updating fragments like infinite fire propagation speed in updating direction.
- For the same reason, we need a differentiation between newly ignited (state 2) and burning (state 3) trees.
- The Algorithm is quite slow because of all the loops and if-statements.

5.2 Random Based Updating

In a random based updating implementation, at the beginning of every time step, one random cell is chosen. It then checks for the rules of the model and acts accordingly. In order to see fires, we need to have instantaneous burning, which is actually permitted by the model. However, it changes the dynamics of the system drastically. But more about that later. The following steps are taken for each timestep:

1. Choose a single cell
2. If cell is empty, grow a tree with probability p
3. If cell is a tree, with probability f , burn it and the connected cluster

In Matlab, this looks something like this:

```
for i=1:t % Loop over all timesteps
(j,k)=ceil(Grid_Size*rand(1,2)); % Choose a random cell
if Grid(j,k)==0 % If the cell is empty
\\With some probability p set G(j,k)==1;
end
if Grid(j,k)==1 % If the cell is alive
\\with some probability f execute burn(j,k);
end
end
```

This implementation of the forest fire model has several advantages:

- There are no problems due to updating processes
- Much shorter implementation
- Easier to extract data

The last point is to be explained more extensively: Every time a tree ignites, the whole connected cluster is burned down. Since the whole process happens in a single timestep, there are no big issues.

5.2.1 keeping track of the clusters

Keeping track of the different clusters was more difficult than we thought at first, but we managed to find an elegant solution. It works by keeping a second grid which stores all the cluster indices. This has several advantages:

- if a new tree grows, we can just look at the neighbor cluster index to determine the current cell index
- if a tree burns, we simply delete all the trees with the same index as the burned tree.

Problems and Solutions Problems arose quickly. Here is a short summary of all the major problems and the corresponding solutions:

Storing the indices The problem here was to find a way to store all possible cluster indices and their availability. Since the maximum number of clusters can be easily found to be $N_{c_{max}} = \frac{G^2}{2}$ (imagine a perfect chessboard-layout filling half of the grid; if one more is added, they combine into a bigger cluster), we already knew the size of the vector. The final format was a Matrix with $N_{c_{max}}$ Columns and 2 rows. The first row would simply store all the possible indices from 1 to $N_{c_{max}}$ and the second row would store the corresponding binary information about the availability.

growing a tree between two clusters It is simple to determine the cluster index of a new tree if there is just one neighbor. The big problem arises if there are more than two neighbors with *different* cluster indices. The solution to this problem was that we set the index as the lower of both neighbors and changed the whole cluster with the higher index to the lower one. This was relatively easy to do by using the matlab built-in function —ismember— which searches for all entries with a certain value and returns a matrix containing only that information.

5.3 The *neighbor* function

The *neighbor* function is the same for both implementations. It takes the current cell coordinates (j, k) and returns a vector containing the coordinates of its direct neighbors. It works like this:

```
Set all neighbors to standard
// Ex: West_X=j-1; West_Y=k;
Treat Special cases
// Ex: if j==1
// West_X=size(Grid,1);
```

Note that the *neighbor* function creates periodic boundary conditions.

5.4 Saving the data

Because we were doing a lot of parameter sweeps and thus generating a lot of data, it was crucial to implement an automatic data saving procedure. This would store all the relevant information of a single simulation in a custom matlab struct and save the important figures. We also created a custom naming concept which allows for quick search for the desired data. All simulation results are named after their input parameters:

6 Simulation Results and Discussion

note that in this section, Θ is defined as p/f .

6.1 Tree population

A first step in the analysis of the results is to have a close look at the tree population during the simulation. Here are two examples:




Figure 1: $N = 50$, $\Theta = 1000$, Simulation time 500'000




Figure 2: $N = 100$, $\Theta = 100$, Simulation time 500'000

What can we learn from these plots? Well, in the first figure, Theta is equal to 10^3 , which is a moderately large number. Yet, the plot of the tree fraction looks like random noise. We are not quite sure what we should interpret from such a plot. At $\Theta = 100$, the plot looks a lot more stable, more like an actual *equilibrium* point. Re-running the simulation with the same Θ but changed f yields the same, consistent picture; after a short transient period to overcome the initial condition, the plot approximates a straight line. Now we know that the SOC state is not generally an equilibrium state and further insight into the properties of the system can only be achieved by analysis of other properties.

6.2 Cluster Radius

The Cluster Radius is defined as the root mean of all the distances of all cells of a cluster to their common center. This property should, given an SOC behaviour of the system, exhibit a power-law-behaviour that would be described as $S(r) = r^{-\beta}$ where $S(r)$ is the number of occurrences of clusters with radius r . In theory, what this means is, if you plot a double-logarithmic histogram of all cluster radii, it should give us a straight line with slope β . Of course, we expected a little "noise", given that our simulation was not infinite in time and space, but it should have approached it pretty well. Here is one of those double-logarithmic histograms:




Figure 3: Cluster radius distribution ($N = 100$, $\Theta = 1000$, Sim Time 500'000)

The Problem that we see here is one that's closely related to the periodic boundary conditions. Imagine an empty grid with a small cluster that stretches from one side *over the boundary* to the other side. Its center will be in the middle, but its radius will be much larger than it actually is. We guess that this causes the tilt that one can observe at medium sized clusters. We are currently working on a way to get around that problem.

6.3 Cluster Size

The cluster size is another variable that we can easily measure. Here is a log-log-histogram of a simulation:




Figure 4: Cluster size distribution ($N = 100$, $\Theta = 1000$, Sim Time 500'000)

7 Open Questions

We are wondering if it is "correct" to implement the forest fire in such a way with zero fire burning time like stated in [1]. We have thus implemented two separate models, one of which uses finite burning speeds, while the other one (described in this text) uses infinite burning speed. Looking at the different timescales, it is obvious that the spreading of a fire should be very fast, but is it correct to use this in the limit of instantaneous fires?

8 Summary and Outlook

References

- [1] Gunnar Pruessner, Henrik Jeldtoft Jensen. *A new, efficient algorithm for the Forest Fire Model*, Phys. Rev. E 70 066707-1–25, 2004
- [2] Per Bak, Chao Tang, Kurt Wiesenfeld. *Self-organized criticality* Phys Rev. A 38 1, 1988